

Reinforcement Learning

Lecture 4: Mathematical Foundations - MDPs and Bellman Equations

Taehoon Kim

Sogang University MIMIC Lab
<https://mimic-lab.com>

Fall Semester 2025

Learning Objectives

By the end of this lecture, you will:

- 1 **Understand** the mathematical formulation of MDPs
- 2 **Master** value functions and their recursive relationships
- 3 **Derive** and apply Bellman equations
- 4 **Implement** policy evaluation and improvement
- 5 **Code** complete Policy Iteration and Value Iteration
- 6 **Analyze** convergence properties and error bounds

Prerequisites:

- Linear algebra (matrix operations)
- Probability theory basics
- PyTorch tensor operations
- Python programming experience

What is a Markov Decision Process?

An MDP is a mathematical framework for sequential decision-making


Components: $\mathcal{M} = (\mathcal{S}, \mathcal{A}, P, r, \gamma)$

- \mathcal{S} : State space (finite or infinite)
- \mathcal{A} : Action space (finite or infinite)
- P : Transition dynamics $P(s'|s, a)$
- r : Reward function $r(s, a) = \mathbb{E}[R_{t+1}|S_t = s, A_t = a]$
- γ : Discount factor $\gamma \in [0, 1]$

Key Property: Markov Property

$$P(S_{t+1}|S_t, A_t, S_{t-1}, A_{t-1}, \dots) = P(S_{t+1}|S_t, A_t) \quad (1)$$

Agent-Environment Interaction



Trajectory: $S_0, A_0, R_1, S_1, A_1, R_2, S_2, \dots$

Example: GridWorld MDP

GridWorld Environment:

- States: Grid cells
- Actions: {UP, RIGHT, DOWN, LEFT}
- Transitions: Move to adjacent cell
- Rewards: -0.04 per step, +1 at goal, -1 at pit
- Discount: $\gamma = 0.99$

S	.	.	G(+1)
.	#	.	P(-1)
.	.	.	.

S = Start, G = Goal, P = Pit, # = Wall

Deterministic vs Stochastic Transitions

Deterministic:

- Action always leads to same next state
- $P(s'|s, a) \in \{0, 1\}$
- Simpler to analyze

Example with slip probability 0.2:

- Intended direction: probability 0.6
- Perpendicular directions: probability 0.2 each

Stochastic:

- Action may lead to different states
- $P(s'|s, a) \in [0, 1]$
- More realistic

A policy π defines the agent's behavior

Types of Policies:

- **Deterministic:** $\pi : \mathcal{S} \rightarrow \mathcal{A}$
 - Maps each state to a single action
 - $a = \pi(s)$
- **Stochastic:** $\pi : \mathcal{S} \times \mathcal{A} \rightarrow [0, 1]$
 - Probability distribution over actions
 - $\pi(a|s) = P(A_t = a | S_t = s)$

Goal: Find the optimal policy π^* that maximizes expected return

Returns and Episodes

Return: Sum of (discounted) future rewards

Finite Horizon Return:

$$G_t = R_{t+1} + R_{t+2} + \dots + R_T \quad (2)$$

Infinite Horizon Discounted Return:

$$G_t = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \quad (3)$$

Why discount?

- Mathematical convergence
- Uncertainty about the future
- Economic interpretation (interest rates)
- Bounded returns: $|G_t| \leq \frac{R_{max}}{1-\gamma}$

State-Value Function

Value of a state under policy π :

$$v_{\pi}(s) = \mathbb{E}_{\pi} \left[\sum_{t=0}^{\infty} \gamma^t R_{t+1} \mid S_0 = s \right] \quad (4)$$

Interpretation:

- Expected return starting from state s
- Following policy π thereafter
- Accounts for all future rewards (discounted)

Properties:

- Bounded: $|v_{\pi}(s)| \leq \frac{R_{max}}{1-\gamma}$
- Unique for a given policy and MDP

Action-Value Function (Q-Function)

Value of taking action a in state s under policy π :

$$q_{\pi}(s, a) = \mathbb{E}_{\pi} \left[\sum_{t=0}^{\infty} \gamma^t R_{t+1} \mid S_0 = s, A_0 = a \right] \quad (5)$$

Relationship to state-value:

$$v_{\pi}(s) = \sum_a \pi(a|s) q_{\pi}(s, a) \quad (6)$$

Q-function tells us:

- How good is action a in state s ?
- Basis for action selection
- Central to Q-learning algorithms

Value Function Examples

Random Policy Values:

0.52	0.55	0.61	+1
0.48	#	0.43	-1
0.44	0.41	0.38	0.35

Optimal Policy Values:

0.81	0.87	0.92	+1
0.76	#	0.66	-1
0.71	0.66	0.61	0.39

Observation: Optimal values are higher (better policy)

Bellman Expectation Equation for v_π

Recursive decomposition of value:

$$v_\pi(s) = \mathbb{E}_\pi[R_{t+1} + \gamma v_\pi(S_{t+1}) | S_t = s] \quad (7)$$

Expanded form:

$$v_\pi(s) = \sum_a \pi(a|s) \sum_{s'} P(s'|s, a) [r(s, a) + \gamma v_\pi(s')] \quad (8)$$

Key insight:

- Value = immediate reward + discounted future value
- Self-consistent system of equations
- One equation per state

Bellman Expectation Equation for q_π

Recursive decomposition:

$$q_\pi(s, a) = r(s, a) + \gamma \sum_{s'} P(s'|s, a) \sum_{a'} \pi(a'|s') q_\pi(s', a') \quad (9)$$

Relationship between V and Q:

$$v_\pi(s) = \sum_a \pi(a|s) q_\pi(s, a) \quad (10)$$

$$q_\pi(s, a) = r(s, a) + \gamma \sum_{s'} P(s'|s, a) v_\pi(s') \quad (11)$$

These equations form the basis for policy evaluation

Matrix Form of Bellman Expectation

For deterministic policy π :

Define:

- $v_\pi \in \mathbb{R}^{|\mathcal{S}|}$: value vector
- $r_\pi \in \mathbb{R}^{|\mathcal{S}|}$: reward vector
- $P_\pi \in \mathbb{R}^{|\mathcal{S}| \times |\mathcal{S}|}$: transition matrix

Bellman equation in matrix form:

$$v_\pi = r_\pi + \gamma P_\pi v_\pi \quad (12)$$

Solution:

$$v_\pi = (I - \gamma P_\pi)^{-1} r_\pi \quad (13)$$

Note: Direct inversion is $O(n^3)$ - iterative methods preferred

Bellman Expectation Operator

Define operator $T_\pi : \mathbb{R}^{|S|} \rightarrow \mathbb{R}^{|S|}$:

$$(T_\pi v)(s) = r_\pi(s) + \gamma \sum_{s'} P_\pi(s'|s) v(s') \quad (14)$$

Properties:

- T_π is a γ -contraction in $\|\cdot\|_\infty$
- $\|T_\pi v - T_\pi w\|_\infty \leq \gamma \|v - w\|_\infty$
- Has unique fixed point v_π
- $v_\pi = T_\pi v_\pi$ (Bellman equation)

Banach Fixed-Point Theorem:

- Starting from any v_0
- Sequence $v_{k+1} = T_\pi v_k$ converges to v_π
- Convergence rate: geometric with factor γ

Optimal Value Functions

Optimal state-value function:

$$v_*(s) = \max_{\pi} v_{\pi}(s) \quad (15)$$

Optimal action-value function:

$$q_*(s, a) = \max_{\pi} q_{\pi}(s, a) \quad (16)$$

Properties:

- Unique for a given MDP
- Defines the best possible performance
- Independent of initial state distribution

Optimal policy: Any policy achieving v_* is optimal

Bellman Optimality Equation for v_*

The optimal value satisfies:

$$v_*(s) = \max_a \left\{ r(s, a) + \gamma \sum_{s'} P(s'|s, a) v_*(s') \right\} \quad (17)$$

Key difference from expectation equation:

- Max over actions instead of expectation
- Non-linear due to max operator
- Harder to solve than expectation equation

Interpretation:

- Optimal value = best immediate reward + discounted future
- Greedy action selection

Bellman Optimality Equation for q_*

The optimal Q-function satisfies:

$$q_*(s, a) = r(s, a) + \gamma \sum_{s'} P(s'|s, a) \max_{a'} q_*(s', a') \quad (18)$$

Relationship between v_* and q_* :

$$v_*(s) = \max_a q_*(s, a) \quad (19)$$

$$q_*(s, a) = r(s, a) + \gamma \sum_{s'} P(s'|s, a) v_*(s') \quad (20)$$

Optimal policy extraction:

$$\pi_*(s) \in \arg \max_a q_*(s, a) \quad (21)$$

Bellman Optimality Operator

Define operator $T_* : \mathbb{R}^{|S|} \rightarrow \mathbb{R}^{|S|}$:

$$(T_* v)(s) = \max_a \left\{ r(s, a) + \gamma \sum_{s'} P(s'|s, a) v(s') \right\} \quad (22)$$

Properties:

- T_* is a γ -contraction
- Unique fixed point v_*
- Monotonic: if $v \leq w$ then $T_* v \leq T_* w$
- Non-linear due to max

Value Iteration: $v_{k+1} = T_* v_k$ converges to v_*

Dynamic Programming for MDPs

Requirements:

- Complete model of environment (P and r known)
- Finite state and action spaces
- Computational resources

Two main algorithms:

- 1 **Policy Iteration:** Alternates evaluation and improvement
- 2 **Value Iteration:** Directly applies optimality operator

Both algorithms:

- Converge to optimal policy
- Based on Bellman equations
- Exact solutions for finite MDPs

Policy Evaluation

Given: Policy π , MDP model

Find: v_π

Algorithm 1 Iterative Policy Evaluation

```
1: Initialize  $V(s) = 0$  for all  $s \in \mathcal{S}$ 
2: repeat
3:    $\Delta \leftarrow 0$ 
4:   for each  $s \in \mathcal{S}$  do
5:      $v \leftarrow V(s)$ 
6:      $V(s) \leftarrow \sum_a \pi(a|s) \sum_{s'} P(s'|s, a) [r(s, a) + \gamma V(s')]$ 
7:      $\Delta \leftarrow \max(\Delta, |v - V(s)|)$ 
8:   end for
9: until  $\Delta < \epsilon$ 
10: return  $V$ 
```

Policy Improvement

Given: Value function v_π

Find: Better policy π'

Greedy policy with respect to v_π :

$$\pi'(s) = \arg \max_a \left\{ r(s, a) + \gamma \sum_{s'} P(s'|s, a) v_\pi(s') \right\} \quad (23)$$

Policy Improvement Theorem:

- If π' is greedy w.r.t. v_π
- Then $v_{\pi'}(s) \geq v_\pi(s)$ for all s
- Equality holds iff π is optimal

This guarantees monotonic improvement!

Algorithm 2 Policy Iteration

```
1: Initialize  $\pi$  arbitrarily
2: repeat
3:   Policy Evaluation:
4:     Compute  $v_\pi$  using iterative evaluation
5:   Policy Improvement:
6:      $\pi_{old} \leftarrow \pi$ 
7:     for each  $s \in \mathcal{S}$  do
8:        $\pi(s) \leftarrow \arg \max_a q_\pi(s, a)$ 
9:     end for
10: until  $\pi = \pi_{old}$ 
11: return  $\pi$ 
```

Convergence: Finite number of iterations for finite MDPs

Algorithm 3 Value Iteration

```
1: Initialize  $V(s) = 0$  for all  $s \in \mathcal{S}$ 
2: repeat
3:    $\Delta \leftarrow 0$ 
4:   for each  $s \in \mathcal{S}$  do
5:      $v \leftarrow V(s)$ 
6:      $V(s) \leftarrow \max_a \sum_{s'} P(s'|s, a)[r(s, a) + \gamma V(s')]$ 
7:      $\Delta \leftarrow \max(\Delta, |v - V(s)|)$ 
8:   end for
9: until  $\Delta < \epsilon$ 
10: Extract policy:  $\pi(s) = \arg \max_a q(s, a)$ 
11: return  $\pi$ 
```

Combines evaluation and improvement in single update

Policy Iteration vs Value Iteration

Aspect	Policy Iteration	Value Iteration
Updates	Alternating	Combined
Convergence	Fewer iterations	More iterations
Per iteration cost	Higher (full evaluation)	Lower (single backup)
Memory	Store policy + values	Store values only
Early stopping	Natural (policy stable)	Needs error bound
Implementation	More complex	Simpler

In practice:

- VI often preferred for simplicity
- PI can be faster for some problems
- Modified PI: partial evaluation

Contraction Mapping Theorem

Theorem: If T is a γ -contraction on complete metric space, then:

- 1 T has unique fixed point v^*
- 2 For any v_0 , sequence $v_{k+1} = Tv_k$ converges to v^*
- 3 $\|v_k - v^*\|_\infty \leq \gamma^k \|v_0 - v^*\|_\infty$

Application to RL:

- Both T_π and T_* are contractions
- Guarantees convergence of DP algorithms
- Provides convergence rate

Convergence rate:

- Geometric with factor γ
- Higher $\gamma \rightarrow$ slower convergence
- Number of iterations $\propto \frac{1}{1-\gamma}$

Error Bounds for Value Iteration

After k iterations of value iteration:

$$\|V_k - v_*\|_\infty \leq \gamma^k \|V_0 - v_*\|_\infty \quad (24)$$

Using Bellman residual $\delta_k = \|V_{k+1} - V_k\|_\infty$:

$$\|V_k - v_*\|_\infty \leq \frac{\gamma}{1 - \gamma} \delta_k \quad (25)$$

Stopping criterion for ϵ -optimal solution:

- Want: $\|V_k - v_*\|_\infty \leq \epsilon$
- Stop when: $\delta_k < \frac{(1-\gamma)\epsilon}{2\gamma}$

These bounds are tight and computable!

Computational Complexity

Per iteration complexity:

- Policy Evaluation: $O(|\mathcal{S}|^2|\mathcal{A}|)$ per iteration
- Policy Improvement: $O(|\mathcal{S}|^2|\mathcal{A}|)$
- Value Iteration: $O(|\mathcal{S}|^2|\mathcal{A}|)$ per iteration

Number of iterations:

- Policy Iteration: $O(|\mathcal{A}|^{|\mathcal{S}|})$ worst case (rarely reached)
- Value Iteration: $O(\frac{1}{1-\gamma} \log \frac{1}{\epsilon})$

Space complexity:

- Transition matrix: $O(|\mathcal{S}|^2|\mathcal{A}|)$
- Value function: $O(|\mathcal{S}|)$
- Policy: $O(|\mathcal{S}|)$

Implementation: Setup

```
1 import torch
2 import numpy as np
3
4 def setup_seed(seed=42):
5     torch.manual_seed(seed)
6     np.random.seed(seed)
7     if torch.cuda.is_available():
8         torch.cuda.manual_seed_all(seed)
9
10 device = torch.device(
11     'cuda' if torch.cuda.is_available()
12     else 'mps' if torch.backends.mps.is_available()
13     else 'cpu'
14 )
15
16 # MDP components shapes
17 # P: [S, A, S] - transition probabilities
18 # R: [S, A] - rewards
19 # V: [S] - values
20 # Q: [S, A] - Q-values
21 # pi: [S] - deterministic policy
```

GridWorld MDP Implementation

```
1 class GridWorldMDP:
2     def __init__(self, grid, terminal_rewards,
3                 step_cost=-0.04, slip_prob=0.1):
4         self.height = len(grid)
5         self.width = len(grid[0])
6
7         # Build state mapping
8         self.state_to_pos = {}
9         self.pos_to_state = {}
10        state_idx = 0
11        for r in range(self.height):
12            for c in range(self.width):
13                if grid[r][c] != '#': # Not a wall
14                    self.state_to_pos[state_idx] = (r, c)
15                    self.pos_to_state[(r, c)] = state_idx
16                    state_idx += 1
17
18        self.n_states = len(self.state_to_pos)
19        self.n_actions = 4 # UP, RIGHT, DOWN, LEFT
```

Building Transition and Reward Matrices

```
1 def build_dynamics(self):
2     P = torch.zeros((self.n_states, self.n_actions,
3                     self.n_states), device=device)
4     R = torch.full((self.n_states, self.n_actions),
5                   self.step_cost, device=device)
6
7     for s in range(self.n_states):
8         if self.is_terminal[s]:
9             P[s, :, s] = 1.0 # Absorbing state
10            R[s, :] = 0.0
11            continue
12
13        for a in range(self.n_actions):
14            # Stochastic transitions with slip
15            for actual_a, prob in self.get_transitions(a):
16                next_s = self.get_next_state(s, actual_a)
17                P[s, a, next_s] += prob
18                R[s, a] += prob * self.get_reward(next_s)
19
20    return P, R
```

Policy Evaluation in PyTorch

```
1 def policy_evaluation(P, R, pi, gamma=0.99, tol=1e-8):
2     """Evaluate a policy using matrix operations"""
3     S = P.shape[0]
4
5     # Convert policy to transition matrix
6     if pi.dim() == 1: # Deterministic
7         P_pi = P[torch.arange(S), pi] # [S, S]
8         R_pi = R[torch.arange(S), pi] # [S]
9     else: # Stochastic
10        P_pi = torch.einsum('sa,sas->ss', pi, P)
11        R_pi = torch.einsum('sa,sa->s', pi, R)
12
13    V = torch.zeros(S, device=P.device)
14
15    for _ in range(1000):
16        V_new = R_pi + gamma * (P_pi @ V)
17        if torch.max(torch.abs(V_new - V)) < tol:
18            break
19        V = V_new
20
21    return V
```


Policy Improvement in PyTorch

```
1 def policy_improvement(P, R, V, gamma=0.99):
2     """Extract greedy policy from value function"""
3     # Compute Q-values
4     #  $Q(s,a) = R(s,a) + \gamma \sum_{s'} P(s,a,s') * V(s')$ 
5     Q = R + gamma * torch.einsum('sas,s->sa', P, V)
6
7     # Extract greedy policy
8     pi = torch.argmax(Q, dim=1)
9
10    return pi, Q
11
12 def compute_q_values(P, R, V, gamma=0.99):
13     """Compute Q-values from value function"""
14     return R + gamma * torch.einsum('sas,s->sa', P, V)
```

Complete Policy Iteration

```
1 def policy_iteration(P, R, gamma=0.99, tol=1e-8):
2     S, A = P.shape[0], P.shape[1]
3     pi = torch.zeros(S, dtype=torch.long) # Start arbitrary
4
5     for iteration in range(100):
6         # Policy Evaluation
7         V = policy_evaluation(P, R, pi, gamma, tol)
8
9         # Policy Improvement
10        pi_new, Q = policy_improvement(P, R, V, gamma)
11
12        # Check convergence
13        if torch.equal(pi_new, pi):
14            print(f"Converged in {iteration + 1} iterations")
15            break
16
17        pi = pi_new
18
19    return pi, V
```

Value Iteration in PyTorch

```
1 def value_iteration(P, R, gamma=0.99, tol=1e-8):
2     S = P.shape[0]
3     V = torch.zeros(S, device=P.device)
4
5     for iteration in range(1000):
6         V_old = V.clone()
7
8         # Bellman optimality update
9         Q = compute_q_values(P, R, V_old, gamma)
10        V = torch.max(Q, dim=1)[0]
11
12        # Check convergence
13        delta = torch.max(torch.abs(V - V_old))
14        if delta < tol:
15            print(f"Converged in {iteration + 1} iterations")
16            break
17
18        # Extract optimal policy
19        Q_final = compute_q_values(P, R, V, gamma)
20        pi = torch.argmax(Q_final, dim=1)
21
22    return pi, V
```

Experiment 1: Environment Setup

File: `exp01_setup.py`

Objectives:

- Verify PyTorch installation
- Test device selection (CUDA/MPS/CPU)
- Initialize MDP tensors
- Verify tensor shapes and operations

Key checks:

- Transition probabilities sum to 1
- Reward bounds are finite
- Bellman operator properties

Expected output: System info, device selection, tensor verification

Experiment 2: GridWorld MDP

File: exp02_gridworld.py

Tasks:

- Build complete GridWorld environment
- Create transition matrix $P[s,a,s']$
- Create reward matrix $R[s,a]$
- Handle walls and terminal states
- Add stochastic transitions (slip)

Verification:

- P matrix: proper probability distribution
- Terminal states are absorbing
- Transitions respect walls

Experiment 3: Policy Evaluation

File: `exp03_policy_evaluation.py`

Implement and test:

- Iterative policy evaluation
- Contraction property verification
- Compare different policies (random, fixed)
- Analyze convergence rates with different γ

Key observations:

- Geometric convergence
- Higher $\gamma \rightarrow$ slower convergence
- Unique fixed point

Experiment 4: Policy Improvement

File: `exp04_policy_improvement.py`

Tasks:

- Compute Q-values from V
- Extract greedy policy
- Verify improvement theorem
- Compare deterministic vs soft improvement

Verification:

- $v_{\pi'} \geq v_{\pi}$ for all states
- Monotonic improvement
- Convergence to optimal

Experiment 5: Policy Iteration

File: `exp05_policy_iteration.py`

Complete implementation:

- Full policy iteration algorithm
- Track policy evolution
- Compare different initial policies
- Visualize convergence

Analysis:

- Number of iterations to convergence
- Policy changes per iteration
- Final optimal policy

Experiment 6: Value Iteration

File: exp06_value_iteration.py

Implement:

- Value iteration algorithm
- Compare with policy iteration
- Test different initializations
- Analyze Bellman optimality operator

Comparisons:

- VI vs PI convergence speed
- Computational efficiency
- Memory requirements

Experiment 7: Stopping Criteria

File: exp07_stopping_criteria.py

Explore:

- Different stopping criteria
- Error bounds computation
- Trade-offs: accuracy vs computation
- Early stopping strategies

Key results:

- Theoretical bounds are tight
- Policy often converges before values
- Span seminorm alternative

Experiment 8: Algorithmic Optimizations

File: `exp08_optimizations.py`

Compare:

- Synchronous vs asynchronous updates
- Prioritized sweeping
- GPU acceleration
- Memory efficiency

Performance metrics:

- Iterations to convergence
- Wall-clock time
- Memory usage
- Scalability

Experiment 9: Integrated Test

File: exp09_integrated_test.py

Complete pipeline:

- Multiple test scenarios
- All algorithms comparison
- Comprehensive visualization
- Performance benchmarking

Deliverables:

- Optimal policies for each scenario
- Algorithm comparison report
- Convergence verification
- Reproducibility guarantee

Key Takeaways

Theory:

- MDPs provide mathematical framework for sequential decisions
- Value functions satisfy Bellman equations
- Bellman operators are contractions \rightarrow unique solutions
- DP algorithms solve MDPs exactly

Practice:

- Implemented complete GridWorld MDP
- Policy evaluation converges geometrically
- Policy improvement guarantees monotonic improvement
- PI and VI converge to same optimal policy

Insights:

- Model-based methods are sample efficient
- Computation scales with state/action space
- Foundation for model-free methods

Mathematical Foundations Summary

Concept	Key Equation
State Value	$v_{\pi}(s) = \mathbb{E}_{\pi}[G_t S_t = s]$
Action Value	$q_{\pi}(s, a) = \mathbb{E}_{\pi}[G_t S_t = s, A_t = a]$
Bellman Expectation	$v_{\pi} = r_{\pi} + \gamma P_{\pi} v_{\pi}$
Bellman Optimality	$v_{*} = \max_a \{r_a + \gamma P_a v_{*}\}$
Policy Evaluation	$v_{k+1} = T_{\pi} v_k$
Value Iteration	$v_{k+1} = T_{*} v_k$
Greedy Policy	$\pi(s) = \arg \max_a q(s, a)$

Remember: These equations are the heart of RL!