

# Reinforcement Learning

## Lecture 8: Policy Gradient Methods - REINFORCE

Taehoon Kim

Sogang University MIMIC Lab  
<https://mimic-lab.com>

Fall Semester 2025

# Today's Agenda

- Motivate policy gradients and contrast them with value-based RL
- Derive the policy gradient theorem and reinforce update
- Study variance reduction with reward-to-go and baselines
- Walk through PyTorch implementation patterns and debugging tips
- Review experimental results from nine hands-on scripts

# Learning Objectives

By the end of this lecture, you will be able to:

- ① Derive and interpret the policy gradient theorem
- ② Implement REINFORCE with and without baselines
- ③ Compare moving-average and learned value function baselines
- ④ Analyze variance, entropy regularization, and normalization effects
- ⑤ Build a complete policy gradient agent in PyTorch

## Prerequisites:

- Understanding of MDPs and value functions (Lectures 4-5)
- Experience with neural networks and PyTorch (Lecture 2)
- Familiarity with Q-learning concepts (Lectures 5-7)

# Why Policy Gradients?

## Value-based methods (Q-learning, DQN):

- Learn  $Q(s, a)$ , derive policy:  $\pi(s) = \arg \max_a Q(s, a)$
- Discrete actions only (or discretized)
- Deterministic policies

## Policy-based methods:

- Directly parameterize  $\pi_\theta(a|s)$
- Natural for continuous actions
- Stochastic policies
- Can learn suboptimal stochastic policies

**Key insight:** Optimize expected return directly!

# Policy Parameterization

**Stochastic policy:**  $\pi_\theta(a|s)$  outputs probability distribution

**Discrete actions (Categorical):**

```
1 logits = neural_network(state) # [batch, n_actions]
2 probs = softmax(logits)
3 action = sample_categorical(probs)
```

**Continuous actions (Gaussian):**

```
1 mean = neural_network(state) # [batch, action_dim]
2 std = exp(log_std_parameter) # learnable or fixed
3 action = sample_normal(mean, std)
```

# Policy Gradient Objective

**Goal:** Maximize expected return

$$J(\theta) = \mathbb{E}_{\tau \sim \pi_\theta}[R(\tau)]$$

where trajectory  $\tau = (s_0, a_0, r_0, s_1, a_1, r_1, \dots)$

**Different formulations:**

- Start state value:  $J(\theta) = V^{\pi_\theta}(s_0)$
- Average value:  $J(\theta) = \sum_s d^{\pi_\theta}(s) V^{\pi_\theta}(s)$
- Average reward:  $J(\theta) = \sum_{s,a} d^{\pi_\theta}(s) \pi_\theta(a|s) r(s, a)$

**Challenge:** How to compute  $\nabla_\theta J(\theta)$ ?

# The Gradient Problem

**Why can't we differentiate directly?**

$$J(\theta) = \sum_{\tau} P(\tau|\theta)R(\tau)$$

Taking gradient:

$$\nabla_{\theta} J(\theta) = \sum_{\tau} \nabla_{\theta} P(\tau|\theta)R(\tau)$$

**Problem:**  $P(\tau|\theta)$  depends on environment dynamics!

$$P(\tau|\theta) = p(s_0) \prod_{t=0}^{T-1} \pi_{\theta}(a_t|s_t)p(s_{t+1}|s_t, a_t)$$

We don't know  $p(s_{t+1}|s_t, a_t)$ !

# The Log-Derivative Trick

**Key insight:** Use the identity

$$\nabla_{\theta} P(\tau|\theta) = P(\tau|\theta) \nabla_{\theta} \log P(\tau|\theta)$$

**Proof:**

$$\nabla_{\theta} \log P(\tau|\theta) = \frac{1}{P(\tau|\theta)} \nabla_{\theta} P(\tau|\theta)$$

**Therefore:**

$$\begin{aligned}\nabla_{\theta} J(\theta) &= \sum_{\tau} P(\tau|\theta) \nabla_{\theta} \log P(\tau|\theta) R(\tau) \\ &= \mathbb{E}_{\tau \sim \pi_{\theta}} [\nabla_{\theta} \log P(\tau|\theta) R(\tau)]\end{aligned}$$

Now it's an expectation - we can sample!

# Simplifying the Gradient

**Log probability of trajectory:**

$$\log P(\tau|\theta) = \log p(s_0) + \sum_{t=0}^{T-1} \log \pi_\theta(a_t|s_t) \quad (1)$$

$$+ \sum_{t=0}^{T-1} \log p(s_{t+1}|s_t, a_t) \quad (2)$$

**Taking gradient w.r.t.  $\theta$ :**

$$\nabla_\theta \log P(\tau|\theta) = \sum_{t=0}^{T-1} \nabla_\theta \log \pi_\theta(a_t|s_t)$$

Environment dynamics cancel out!

# Policy Gradient Theorem

## Theorem (Policy Gradient)

*The gradient of expected return is:*

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[ \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) G_t \right]$$

where  $G_t = \sum_{k=t}^{T-1} \gamma^{k-t} r_k$  is the return from time  $t$ .

### Intuition:

- Increase probability of actions that lead to high returns
- Decrease probability of actions that lead to low returns
- Weight by how good the return was

# Score Function Estimator

**Score function:**  $\nabla_{\theta} \log \pi_{\theta}(a|s)$

**For discrete actions (softmax):**

$$\nabla_{\theta} \log \pi_{\theta}(a|s) = \phi(s, a) - \mathbb{E}_{a' \sim \pi}[\phi(s, a')]$$


**For continuous actions (Gaussian):**

$$\nabla_{\theta} \log \mathcal{N}(a|\mu_{\theta}(s), \sigma^2) = \frac{(a - \mu_{\theta}(s))}{\sigma^2} \nabla_{\theta} \mu_{\theta}(s)$$

**Properties:**

- $\mathbb{E}_{a \sim \pi}[\nabla_{\theta} \log \pi_{\theta}(a|s)] = 0$
- Points in direction of increasing action probability

# Experiment 2: Score-Function Visualisation



Exported by `exp02_policy_gradient_math.py`.

## Observations

- Analytical gradient at  $\theta = 0$  is **0.50**; one Monte Carlo draw produced **-0.0004**
- Shows the score-function surface and gradient direction for a two-action policy
- Single-sample estimates can be extremely noisy even in this toy example
- Motivates baselines, reward-to-go, and batching strategies studied later

# REINFORCE: Monte Carlo Policy Gradient

---

**Algorithm 1** REINFORCE (Basic Version)

```
1: Initialize policy network  $\pi_\theta$ 
2: for episode = 1, 2, ... do
3:   Collect trajectory  $\tau = (s_0, a_0, r_0, \dots, s_{T-1}, a_{T-1}, r_{T-1})$ 
4:   for  $t = 0$  to  $T - 1$  do
5:      $G_t \leftarrow \sum_{k=t}^{T-1} \gamma^{k-t} r_k$     (compute return)
6:   end for
7:    $\nabla J \leftarrow \sum_{t=0}^{T-1} \nabla_\theta \log \pi_\theta(a_t | s_t) \cdot G_t$ 
8:    $\theta \leftarrow \theta + \alpha \nabla J$ 
9: end for
```

---

# REINFORCE in PyTorch

```
1 def reinforce_update(policy, optimizer, episode):
2     states, actions, rewards = episode
3
4     # Compute returns
5     returns = []
6     G = 0
7     for r in reversed(rewards):
8         G = r + gamma * G
9         returns.insert(0, G)
10
11    # Compute loss
12    loss = 0
13    for s, a, G in zip(states, actions, returns):
14        logits = policy(s)
15        log_prob = F.log_softmax(logits)[a]
16        loss += -log_prob * G
17
18    # Update
19    optimizer.zero_grad()
20    loss.backward()
21    optimizer.step()
```

# The Variance Problem

**REINFORCE gradient estimator:**

$$\hat{g} = \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) G_t$$

**Properties:**

- **Unbiased:**  $\mathbb{E}[\hat{g}] = \nabla_{\theta} J(\theta)$
- **High variance:** Returns can vary wildly

**Consequences:**

- Slow learning
- Unstable training
- Poor sample efficiency

**Solution:** Variance reduction techniques!

# Variance Reduction: Reward-to-Go

Original REINFORCE:

$$\nabla_{\theta} J = \mathbb{E} \left[ \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) G_0 \right]$$

Reward-to-go:

$$\nabla_{\theta} J = \mathbb{E} \left[ \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) G_t \right]$$

Why does this work?

- Past rewards are independent of future actions
- $\mathbb{E}[\nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \cdot \text{past rewards}] = 0$
- Reduces variance without bias

# Variance Reduction: Baselines

**Baseline subtraction:**

$$\nabla_{\theta} J = \mathbb{E} \left[ \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) (G_t - b(s_t)) \right]$$

**Why unbiased?**

$$\mathbb{E}_{a \sim \pi} [\nabla_{\theta} \log \pi_{\theta}(a | s) \cdot b(s)] = b(s) \cdot \underbrace{\mathbb{E}_{a \sim \pi} [\nabla_{\theta} \log \pi_{\theta}(a | s)]}_{=0} = 0$$

**Common baselines:**

- Constant:  $b = \mathbb{E}[G]$
- State-dependent:  $b(s) = V^{\pi}(s)$  (learned value function)
- Exponential moving average of returns

# Advantage Functions

**Advantage:** How much better is action  $a$  than average?

$$A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s)$$

**Policy gradient with advantages:**

$$\nabla_\theta J = \mathbb{E}_{\tau \sim \pi_\theta} \left[ \sum_{t=0}^{T-1} \nabla_\theta \log \pi_\theta(a_t | s_t) A_t \right]$$

where  $A_t = G_t - V(s_t)$

**Benefits:**

- Much lower variance
- Still unbiased
- Bridge to Actor-Critic methods

# Baseline Comparison

Baseline Type	Complexity	Variance	Bias
None	Low	High	None
Constant	Low	Medium-High	None
EMA	Low	Medium	None
Value Function	High	Low	None*

\*If value function is accurate

## Implementation complexity vs performance trade-off

- Simple baselines: Easy to implement, moderate improvement
- Value function: More complex, best performance
- Path to Actor-Critic methods

# Entropy Regularization

## Modified objective:

$$J(\theta) = \mathbb{E}_{\pi_\theta}[R] + \beta \mathcal{H}(\pi_\theta)$$

where  $\mathcal{H}(\pi_\theta) = -\mathbb{E}_{s,a}[\log \pi_\theta(a|s)]$

## Benefits:

- Encourages exploration
- Prevents premature convergence
- Smoother optimization landscape

## Implementation:

```
1 entropy = -(probs * log_probs).sum(dim=-1)
2 loss = -log_prob * advantage - beta * entropy
```

# Normalization Techniques

## 1. Advantage normalization:

```
1 advantages = (advantages - advantages.mean()) /  
2     (advantages.std() + 1e-8)
```

## 2. Return normalization:

```
1 returns = (returns - returns.mean()) /  
2     (returns.std() + 1e-8)
```

## 3. Observation normalization:

```
1 # Running mean/std  
2 obs_normalized = (obs - running_mean) /  
3     sqrt(running_var + 1e-8)
```

**Benefits:** More stable gradients, faster convergence

# Complete REINFORCE Architecture

## System Components:

- ① **Environment:** Provides states and rewards
- ② **Policy Network**  $\pi_\theta$ : Maps states to action probabilities
- ③ **Value Network**  $V_\phi$ : Estimates state values (for baseline)
- ④ **Episode Buffer:** Stores trajectories  $(s, a, r)$

## Data Flow:

- Environment → State → Policy → Action → Environment
- Episode data → Buffer → Compute returns/advantages
- Gradients update both Policy and Value networks

# Neural Network Design

```
1  class PolicyNetwork(nn.Module):
2      def __init__(self, obs_dim, n_actions):
3          super().__init__()
4          self.net = nn.Sequential(
5              nn.Linear(obs_dim, 128),
6              nn.Tanh(),
7              nn.Linear(128, 128),
8              nn.Tanh(),
9              nn.Linear(128, n_actions)
10         )
11
12     def forward(self, x):
13         return self.net(x)  # logits
14
15     def get_action(self, state):
16         logits = self(state)
17         dist = Categorical(logits=logits)
18         action = dist.sample()
19         return action, dist.log_prob(action)
```

# Training Loop Structure

```
1  for episode in range(num_episodes):
2      # Collect episode
3      states, actions, rewards = [], [], []
4      obs = env.reset()
5
6      while not done:
7          action, log_prob = policy.get_action(obs)
8          next_obs, reward, done, _ = env.step(action)
9          states.append(obs)
10         actions.append(action)
11         rewards.append(reward)
12         obs = next_obs
13
14      # Compute returns and advantages
15      returns = compute_returns(rewards, gamma)
16      advantages = returns - value_net(states)
17
18      # Update networks
19      update_policy(policy, states, actions, advantages)
20      update_value(value_net, states, returns)
```

# Key Hyperparameters

Parameter	Typical Value	Notes
Learning rate (policy)	$10^{-3}$ to $10^{-2}$	Higher than value
Learning rate (value)	$10^{-3}$ to $10^{-4}$	Lower for stability
Discount factor $\gamma$	0.99	Problem-dependent
Entropy coefficient $\beta$	0.01 to 0.001	Decay over time
Gradient clipping	0.5 to 1.0	Prevents explosions
Episodes per update	1 to 16	Trade-off
Hidden dimensions	64 to 256	Task complexity

## Tips:

- Start with high entropy, decay gradually
- Use adaptive optimizers (Adam)
- Monitor gradient norms

# Common Implementation Pitfalls

## 1. Incorrect advantage calculation:

- Use detached values: `advantages = returns - values.detach()`
- Don't backprop through advantages

## 2. Wrong return computation:

- Ensure correct discounting
- Handle episode termination properly

## 3. Gradient issues:

- Always clear gradients: `optimizer.zero_grad()`
- Clip gradients to prevent explosions

## 4. Numerical instability:

- Add small epsilon: `std + 1e-8`
- Use log-sum-exp tricks

# Batch Training

## Single episode update:

- High variance gradients
- Poor GPU utilization
- Unstable learning

## Batch episode update:

```
1 # Collect multiple episodes
2 episodes = [collect_episode() for _ in range(batch_size)]
3
4 # Concatenate all data
5 all_states = concatenate([e.states for e in episodes])
6 all_advantages = concatenate([e.advantages for e in episodes])
7
8 # Single gradient update
9 loss = compute_loss(all_states, all_advantages)
```

Benefits: Lower variance, better GPU usage, more stable

# Learning Rate Scheduling

## Why schedule learning rates?

- Large LR initially for exploration
- Small LR later for convergence
- Adapt to training progress

## Common schedules:

```
1 # Linear decay
2 lr = lr_start * (1 - progress)
3
4 # Exponential decay
5 lr = lr_start * decay_rate ** epoch
6
7 # Cosine annealing
8 lr = lr_min + 0.5 * (lr_max - lr_min) *
9     (1 + cos(pi * epoch / max_epochs))
```

# Gradient Clipping Strategies

## Why clip gradients?

- Prevent gradient explosions
- Stabilize training
- Handle outlier trajectories

## Clipping methods:

```
1 # Clip by value
2 torch.nn.utils.clip_grad_value_(parameters, clip_value)
3
4 # Clip by norm (preferred)
5 torch.nn.utils.clip_grad_norm_(parameters, max_norm)
6
7 # Adaptive clipping
8 if grad_norm > threshold:
9     scale = threshold / grad_norm
10    gradients *= scale
```

# Extension to Continuous Actions

## Gaussian policy:

```
1 class ContinuousPolicy(nn.Module):
2     def __init__(self, obs_dim, action_dim):
3         super().__init__()
4         self.mean_net = nn.Sequential(
5             nn.Linear(obs_dim, 128),
6             nn.Tanh(),
7             nn.Linear(128, action_dim)
8         )
9         self.log_std = nn.Parameter(torch.zeros(action_dim))
10
11    def forward(self, state):
12        mean = self.mean_net(state)
13        std = self.log_std.exp()
14        dist = Normal(mean, std)
15        action = dist.sample()
16        log_prob = dist.log_prob(action).sum(-1)
17        return action, log_prob
```

# From REINFORCE to Actor-Critic

Method	Update	Characteristics
REINFORCE	Monte Carlo	High variance, slow
REINFORCE + baseline	Monte Carlo	Lower variance
Actor-Critic	TD(0)	Online, lower variance
A2C	Synchronous AC	Parallel environments
A3C	Asynchronous AC	Distributed training

**Key difference:** When to update

- REINFORCE: End of episode (Monte Carlo)
- Actor-Critic: Every step (TD learning)

Next lecture: Actor-Critic methods!

# Experiment 3: Vanilla REINFORCE (CartPole-v1)


## Metrics (200 episodes)

- First 10 episodes mean return: **11.7**
- Last 10 episodes mean return: **9.4**
- Greedy evaluation (10 runs):  **$9.2 \pm 1.0$**
- Policy entropy collapses to  $\approx 0$  after  $\sim 30$  episodes
- Loss oscillates between **-0.22** and **0.30** (high variance)

## Takeaways

- Full-return REINFORCE stalls near random performance on CartPole-v1.
- Motivates variance-reduction tricks (reward-to-go, baselines) introduced next.


# Experiment 3: Learning Curve



# Experiment 6: Learned Value Baseline

## CartPole-v1 diagnostic

- Learned  $V(s)$  baseline mean return (last 50):  $9.46 \pm 0.81$
- No-baseline control:  $217.3 \pm 165.9$  (solved by episode 60)
- Value loss drops from 17.9 to 0.36 but policy underfits
- Indicates strong coupling between baseline accuracy and policy LR



## Debug steps

- Increase policy learning rate once baseline stabilizes
- Or bootstrap with TD targets (Actor-Critic, next lecture)
- Monitor value loss + return gap to detect underfitting

# Experiment 8: Advanced REINFORCE Stack

## Configuration

- Batch updates: 8 episodes/update, cosine LR schedule
- Baseline: value network + advantage normalization
- Regularizers: entropy decay  $0.01 \rightarrow 0.001$ , grad clip 0.5

## Outcome

- Basic REINFORCE solved CartPole by update **30**
- Advanced stack underperformed (final **105 ± 53**) — overly heavy regularization
- Highlights need for careful scheduler tuning when batching updates




Figure from `exp08_advanced_techniques.py`: baseline (blue) vs advanced (orange).


# Experiment 9: Integrated REINFORCE Smoke Test

Baseline	Final Eval	Best Eval	Episodes
None	500.0	500.0	400
EMA ( $\alpha = 0.05$ )	500.0	500.0	400
Value	500.0	500.0	400

## Automation checklist

- TensorBoard logs written to `runs/reinforce_*`
- Checkpoints in `checkpoints/`
- Evaluation every 10 updates (10 episodes each)

**Use it to:** sanity-check new baselines, compare seeds, capture regressions



Output from `exp09_integrated_test.py`: training vs evaluation across baselines.


# Experiment 4: Reward-to-Go vs Full Returns

## Return statistics (last 50 episodes)

- Full returns:  $63.9 \pm 43.2$
- Reward-to-go:  $25.6 \pm 4.6$
- Gradient norm (avg): full **0.27**, reward-to-go **0.47**
- Reward-to-go converged slower in this run (needs more tuning)

## Interpretation

- Reward-to-go removes past rewards from the estimator, reducing bias
- With the current hyperparameters it also reduced signal magnitude
- Try pairing reward-to-go with lower learning rate and a baseline



Comparison generated by `exp04_reward_to_go.py`: reward-to-go (blue) vs full returns (orange).


# Experiment 7: Entropy Regularization Sweep

## CartPole-v1 summary

- $\beta = 0$ : final return **9.3**, entropy **0.0000**
- $\beta = 0.001$ : final return **32.8**, entropy **0.060**
- $\beta = 0.01$ : final return **9.5**, entropy **0.0000**
- $\beta = 0.1$ : final return **9.5**, entropy **0.0006**

## Lessons

- A small coefficient ( $10^{-3}$ ) delayed collapse and yielded higher returns
- Large coefficients hurt learning once policy should exploit
- Schedule idea: start at  $10^{-3}$  and decay toward  $10^{-4}$  after stability




Results from `exp07_entropy_regularization.py`.

# Experiment 5: EMA Baseline Sweep

$\alpha$	$\bar{R}_{50}$	$\sigma_{50}$	Episodes $\geq 195$
No baseline	9.26	0.77	0
0.01	<b>195.70</b>	151.11	46 (Ep. 154)
0.05	9.30	0.70	0
0.10	9.32	0.81	0

## Observations

- Small smoothing ( $\alpha = 0.01$ ) occasionally solves CartPole but with high variance
- Aggressive smoothing ( $\alpha \geq 0.05$ ) tracks returns closely  $\rightarrow$  little benefit
- Use EMA only if tuned; pair with reward normalization and lower learning rate



Plot exported by `exp05_baseline_ema.py`: moving baseline vs returns.

# When to Use Policy Gradients?

## Good for:

- Continuous action spaces
- Stochastic optimal policies
- High-dimensional action spaces
- Learning diverse behaviors

## Not ideal for:

- Sample efficiency critical
- Discrete, small action spaces (use DQN)
- Need deterministic policy
- Limited computational budget

**Rule of thumb:** Start with DQN for discrete, PG for continuous

# Debugging Policy Gradients

## Common issues and solutions:

### ① No learning:

- Check gradient flow
- Verify advantage signs
- Increase learning rate

### ② Unstable training:

- Reduce learning rate
- Add gradient clipping
- Normalize advantages

### ③ Policy collapse:

- Add entropy regularization
- Check for numerical issues
- Reduce batch size

# What to Monitor

## Essential metrics:

- Episode returns (mean, std)
- Policy entropy
- Gradient norms
- Value function accuracy (if used)

## TensorBoard logging:

```
1 writer.add_scalar('train/return', episode_return, step)
2 writer.add_scalar('train/entropy', entropy, step)
3 writer.add_scalar('train/grad_norm', grad_norm, step)
4 writer.add_scalar('train/value_loss', value_loss, step)
```

**Warning signs:** Entropy → 0, gradient explosion, value divergence

# Computational Efficiency

## CPU vs GPU:

- Small networks: CPU often sufficient
- Batch processing: GPU advantageous
- Environment step: Always CPU

## Optimization tips:

```
1 # Batch operations
2 states = torch.stack(episode_states)
3 values = value_net(states) # Single forward pass
4
5 # Vectorized environments
6 envs = VectorEnv([make_env() for _ in range(n_envs)])
```

**Memory management:** Clear gradients, detach when needed, limit buffer size

# Key Takeaways

- ➊ **Policy Gradient Theorem:** Direct optimization of expected return
- ➋ **REINFORCE:** Simple but high-variance Monte Carlo method
- ➌ **Variance Reduction:** Critical for practical performance
  - Reward-to-go
  - Baselines (especially value functions)
  - Normalization
- ➍ **Implementation:** Many details matter
  - Proper advantage computation
  - Gradient clipping
  - Entropy regularization
- ➎ **Trade-offs:** Simplicity vs efficiency vs stability

# Algorithm Comparison

Aspect	Q-Learning	DQN	REINFORCE	A2C
Action space	Discrete	Discrete	Any	Any
Update	TD	TD	MC	TD
Sample efficiency	High	High	Low	Medium
Stability	High	Medium	Low	Medium
Continuous	No	No	Yes	Yes
Implementation	Simple	Complex	Simple	Medium

**Next lecture:** Actor-Critic methods (A2C)

- Combines policy gradient with value learning
- Online updates (no waiting for episode end)
- Better sample efficiency

# Next Week: Actor-Critic Methods

## What we'll cover:

- TD learning for value estimation
- Advantage Actor-Critic (A2C)
- Parallel environment collection
- Generalized Advantage Estimation (GAE)

## Key improvements over REINFORCE:

- Online learning (no episode boundary needed)
- Lower variance through bootstrapping
- Better sample efficiency
- Foundation for modern algorithms (PPO, SAC)

**Preparation:** Review TD learning and value functions